

CNN 이란? (Convolutional Neural Network)

KUBIG 12기 이나윤

INDEX

2021년 여름방학 딥러닝 분반 3주차

01

CNN의 개념

02

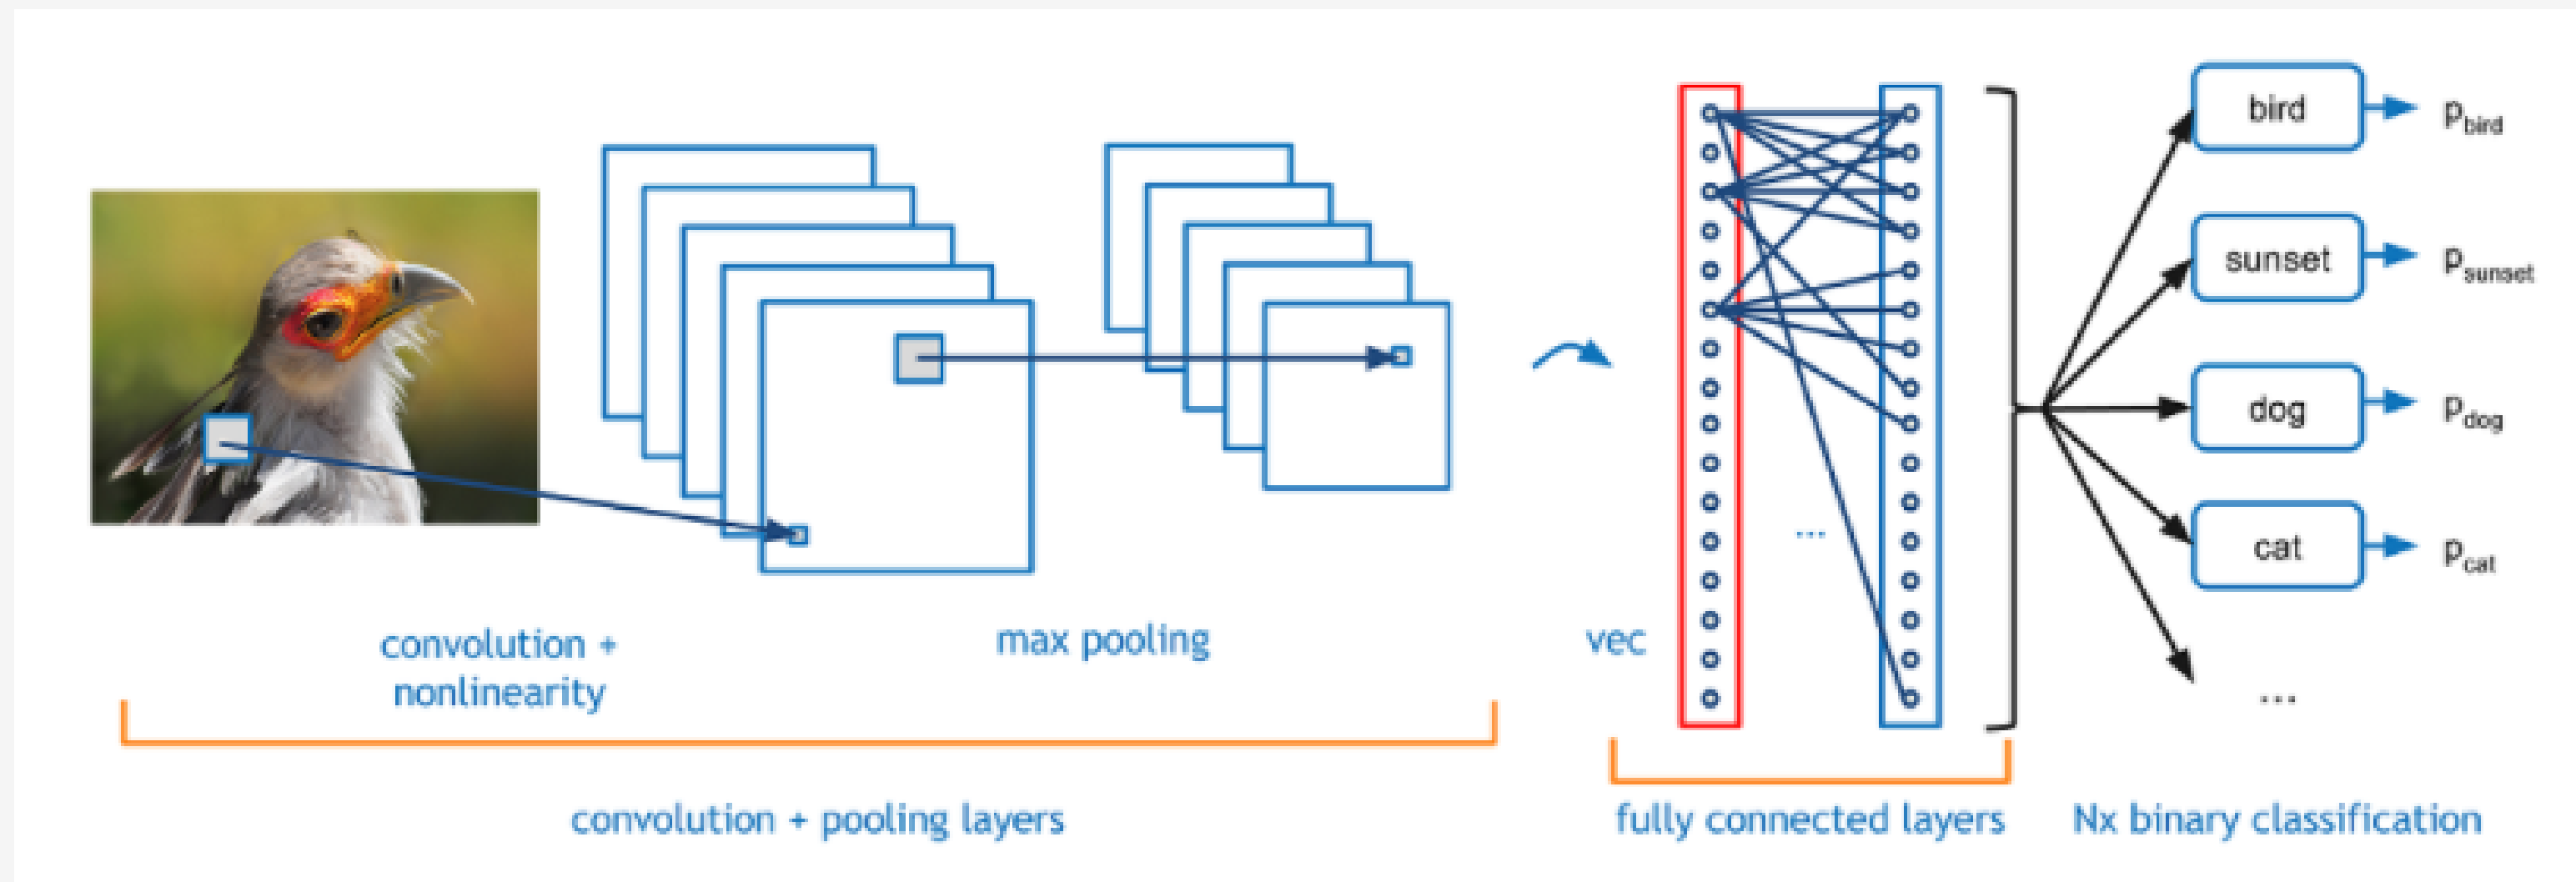
**CNN의 흐름
및 모델**

03

**CNN 관련
논문 리뷰**

01. CNN이란?

CNN(Convolutional Neural Network)




01. CNN – Convolutional Layer/ Fully Connected Layer

CNN(Convolutional Neural Network)


Fully Connected Layer

: 이전 계층의 모든 뉴런이 현재 계층의 **모든 뉴런과 연결된 경우**



Convolutional Layer

: 합성곱으로 이루어진 뉴런을 Fully Connected 형태로 연결한 경우


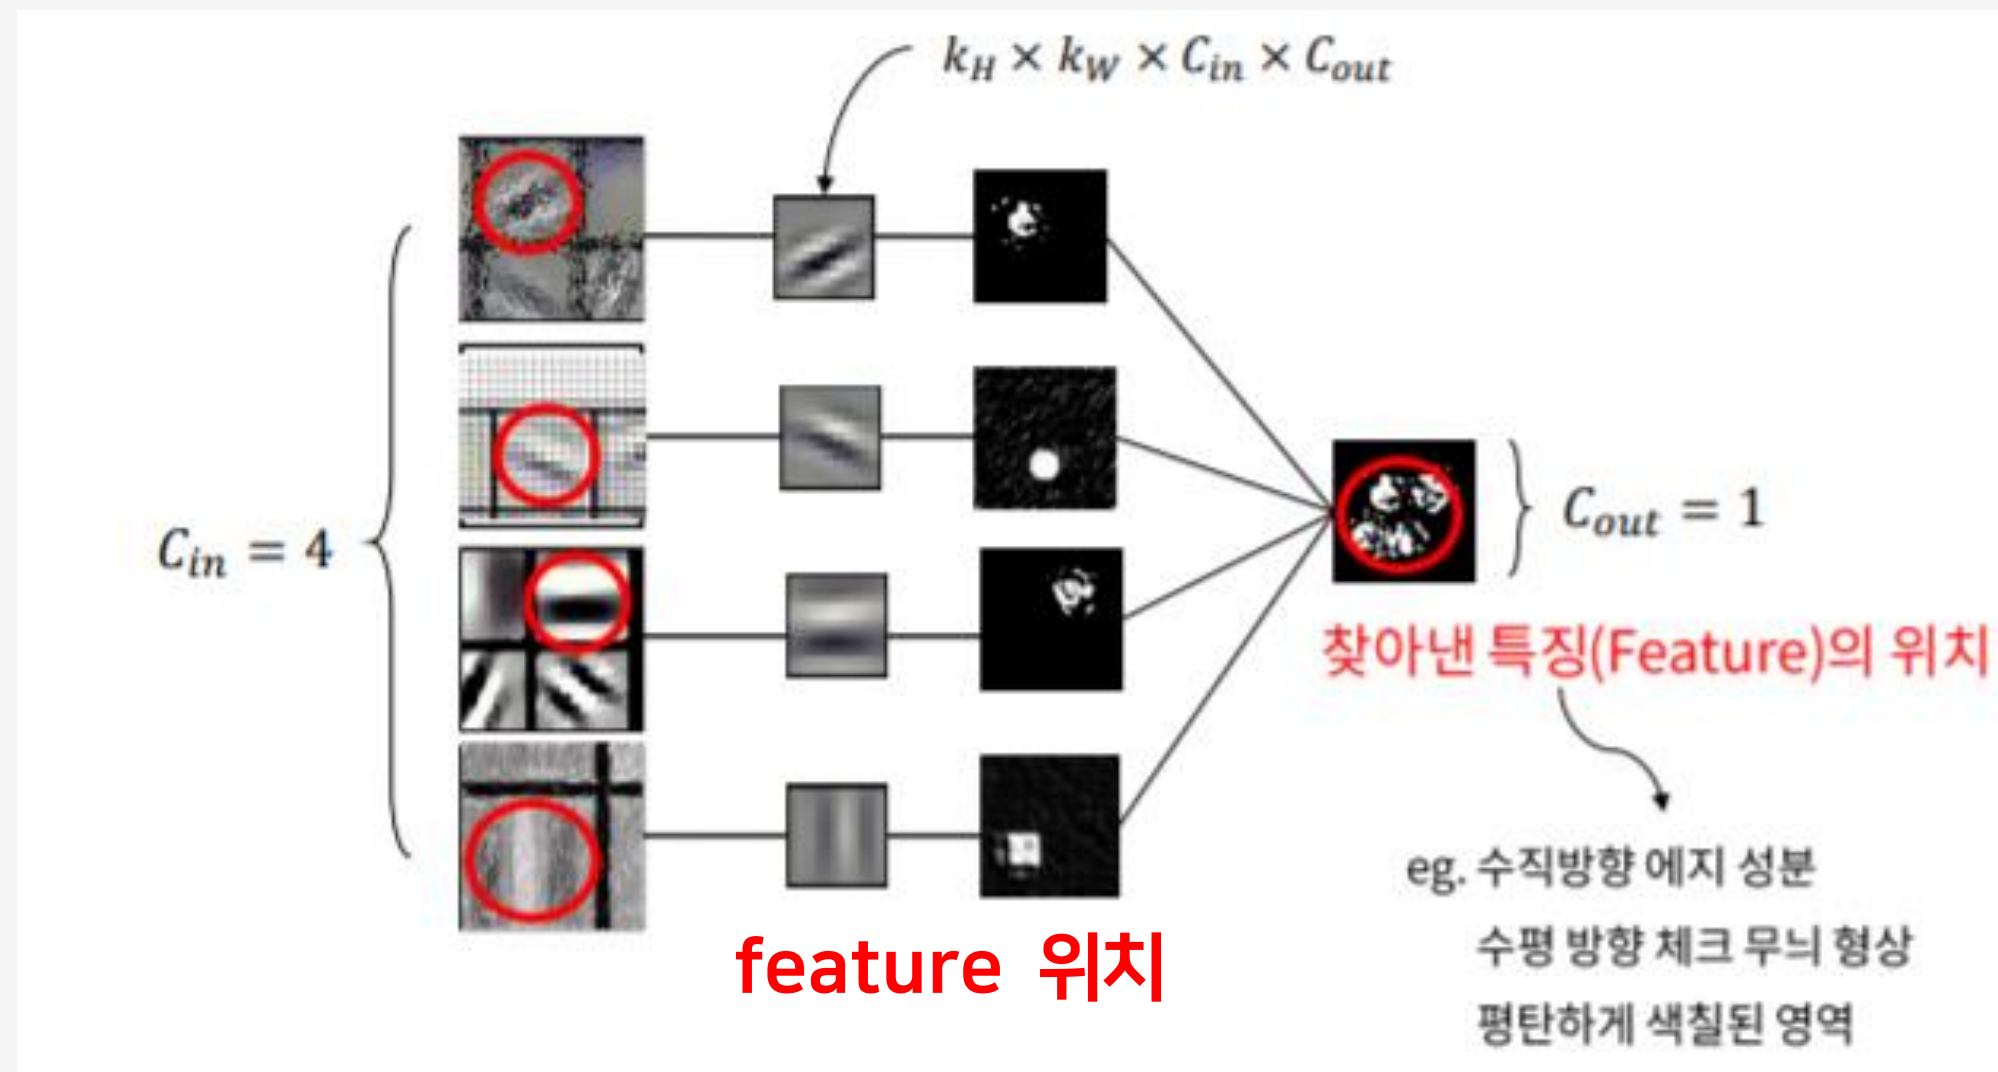


01. CNN – Convolutional Layer


Convolutional Layer

: Consolve the filter with the image

: Slide over the image spatially, computing dot products



01. CNN – Convolutional Layer



Output depth = filter(kernel) 개수


01. CNN – Pooling layer

Pooling Layer : convolution layer의 출력 데이터를 받아, 크기를 줄이거나 특정 데이터를 강조하는 용도로 사용 (정보 종합)

- > 학습대상 Parameter 없음
- > pooling layer를 지나도, 채널 수 변경 없음

Pooling layer

- makes the representations smaller and more manageable
- operates over each activation map independently:



01. CNN – Pooling layer

Pooling Layer

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4



6	8
3	4

Max Pooling

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4




3.25	5.25
2	2

Average Pooling

01. CNN - Stride

Stride: 입력 이미지에서 필터를 몇 칸씩 건너뛰며 적용할지



01. CNN - Padding

padding을 추가하지 않고, filter 통과시킬 경우:

0	3	2	1
2	1	0	2
1	3	2	0
2	0	3	3

 \times

1	0	1
0	1	0
1	0	1

 $=$

6	7
10	8

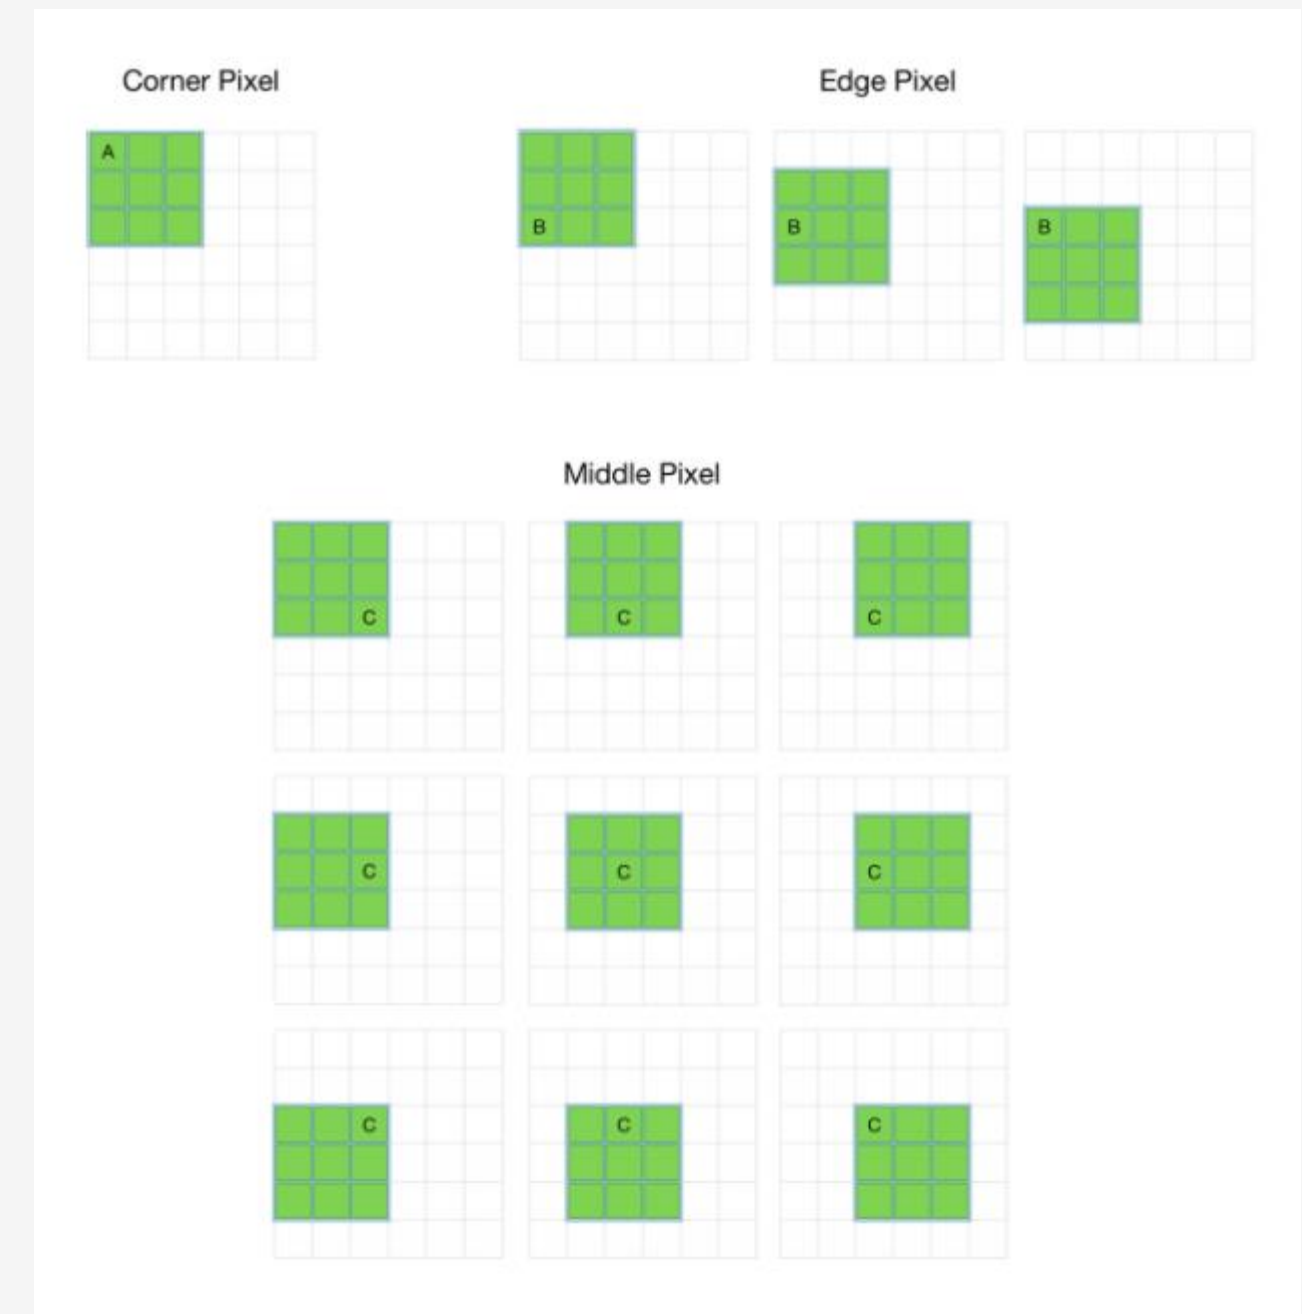
입력 데이터 (4,4) 필터 (3,3) 피쳐맵 (2,2)

1	2	2	1
2	4	4	2
2	4	4	2
1	2	2	1

입력 데이터의
원소별 연산 참여 횟수

-> 이미지 데이터의 축소를 막기 위해

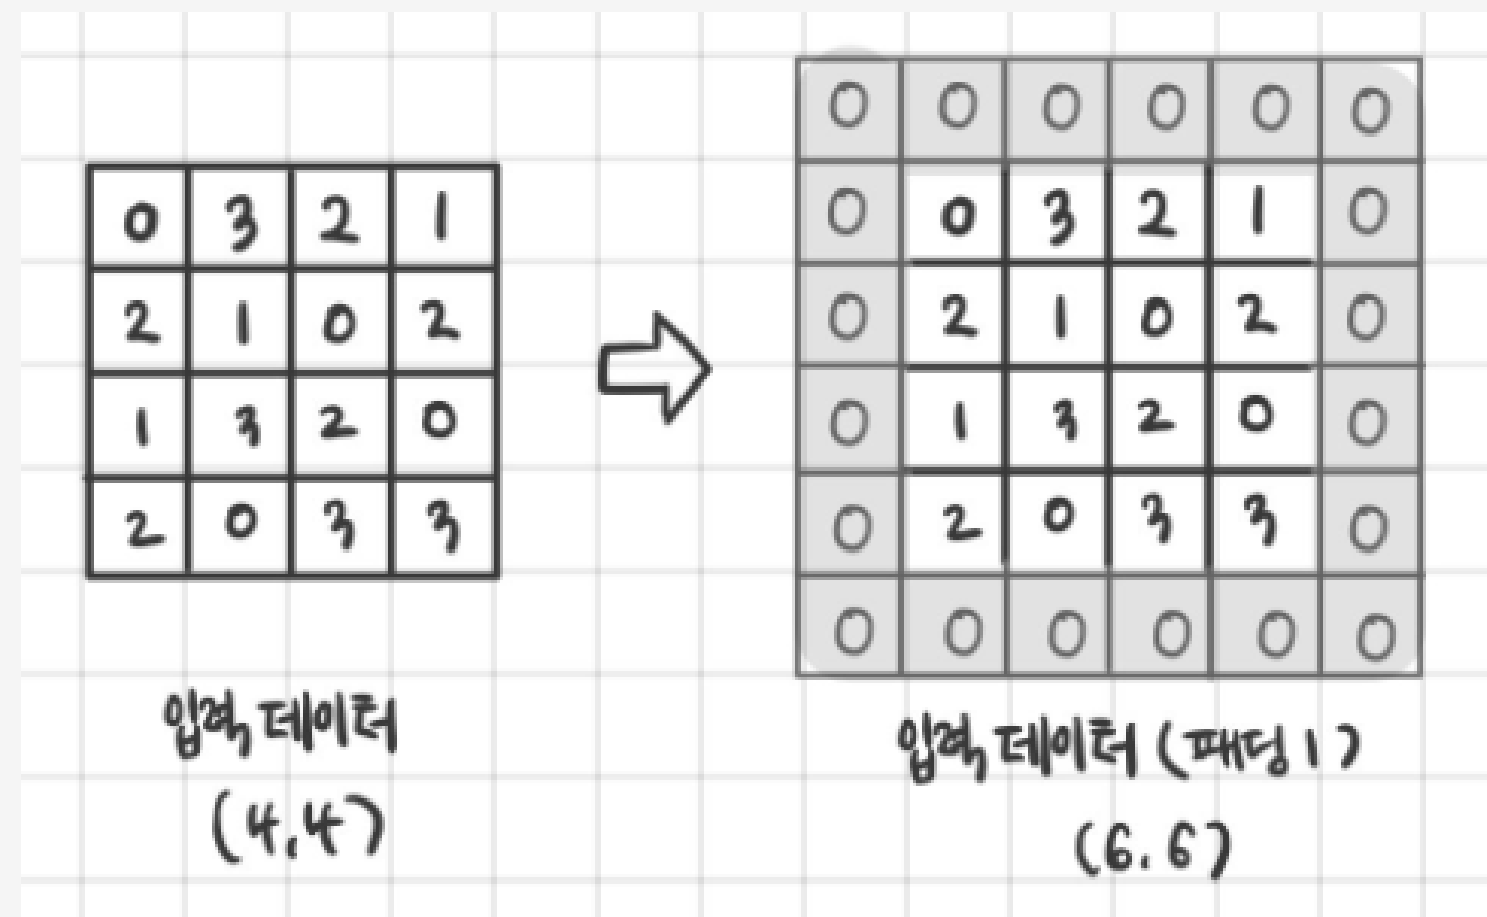
-> Edge data를 충분히 활용하기 위해



01. CNN - Padding

(zero) Padding:

출력 크기를 보정하기 위해 입력 데이터의 사방을 특정 값(일반적으로 0)으로 채우는 것




-> output size = $(n + 2 * p - f) / s + 1$ (bias)

(n,n): origin data, p: padding, f: filter, s: stride

01. CNN이란?

Examples time:

Input volume: **32x32x3**
10 5x5 filters with stride **1**, pad **2**



Output volume size:
 $(32+2*2-5)/1+1 = 32$ spatially, so
32x32x10

Example: CONV layer in Torch

Summary. To summarize, the Conv Layer:

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
 - Number of filters K
 - their spatial extent F ,
 - the stride S ,
 - the amount of zero padding P .

SpatialConvolution

```
module = nn.SpatialConvolution(nInputPlane, nOutputPlane, kW, kH, [dW], [dH], [padW], [padH])
```

Applies a 2D convolution over an input image composed of several input planes. The `input` tensor in `forward(input)` is expected to be a 3D tensor (`nInputPlane` x height x width).

The parameters are the following:

- `nInputPlane` : The number of expected input planes in the image given into `forward()`.
- `nOutputPlane` : The number of output planes the convolution layer will produce. $\rightarrow K$
- `kW` : The kernel width of the convolution
- `kH` : The kernel height of the convolution $\rightarrow P$
- `dW` : The step of the convolution in the width dimension. Default is `1`. $\rightarrow S$
- `dH` : The step of the convolution in the height dimension. Default is `1`.
- `padW` : The additional zeros added per width to the input planes. Default is `0`, a good number is $(kW-1)/2$.
- `padH` : The additional zeros added per height to the input planes. Default is `padW`, a good number is $(kH-1)/2$. $\rightarrow F$

Note that depending of the size of your kernel, several (of the last) columns or rows of the input image might be lost. It is up to the user to add proper padding in images.

If the input image is a 3D tensor `nInputPlane` x height x width , the output image size will be `nOutputPlane` x oheight x owidth where

```
owidth = floor((width + 2*padW - kW) / dW + 1)
oheight = floor((height + 2*padH - kH) / dH + 1)
```

02. CNN의 흐름 및 모델

CNN(Convolutional Neural Network)

: 1989년 'Backpropagation applied to handwritten zip code recognition, Le Cun'

에서 처음 소개됨.

: 이미지를 flat하게 MLP로 학습하면, 지역적 정보가 날아가버림.


따라서, 바로 연산을 하여 지역적 정보를 잘 살릴 수 있는 모델인 CNN이 고안됨.

02. CNN의 흐름 및 모델

1998: LeNet – Gradient-based Learning Applied to Document Recognition

Case Study: LeNet-5


[LeCun et al., 1998]



Conv filters were 5x5, applied at stride 1
Subsampling (Pooling) layers were 2x2 applied at stride 2
i.e. architecture is [CONV-POOL-CONV-POOL-CONV-FC]

02. CNN의 흐름 및 모델

2012: AlexNet – ImageNet Classification with Deep Convolutional Neural Network



Input: 227x227x3 images

First layer (CONV1): 96 11x11 filters applied at stride 4

=>

Q: what is the output volume size? Hint: $(227-11)/4+1 = 55$

Second layer (POOL1): 3x3 filters applied at stride 2

Q: what is the output volume size? Hint: $(55-3)/2+1 = 27$

Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1

[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons

[4096] **FC7**: 4096 neurons

[1000] **FC8**: 1000 neurons (class scores)

Details/Retrospectives:

- first use of ReLU
- used Norm layers (not common anymore)
- heavy data augmentation
- dropout 0.5
- batch size 128
- SGD Momentum 0.9
- Learning rate 1e-2, reduced by 10 manually when val accuracy plateaus
- L2 weight decay 5e-4
- 7 CNN ensemble: 18.2% -> 15.4%

02. CNN의 흐름 및 모델

2014: VggNet – Very Deep Convolutional Networks for Large-Scale Image Recognition

Table 1: **ConvNet configurations** (shown in columns). The depth of the configurations increases from the left (A) to the right (E), as more layers are added (the added layers are shown in bold). The convolutional layer parameters are denoted as “conv(receptive field size)-(number of channels)”. The ReLU activation function is not shown for brevity.

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

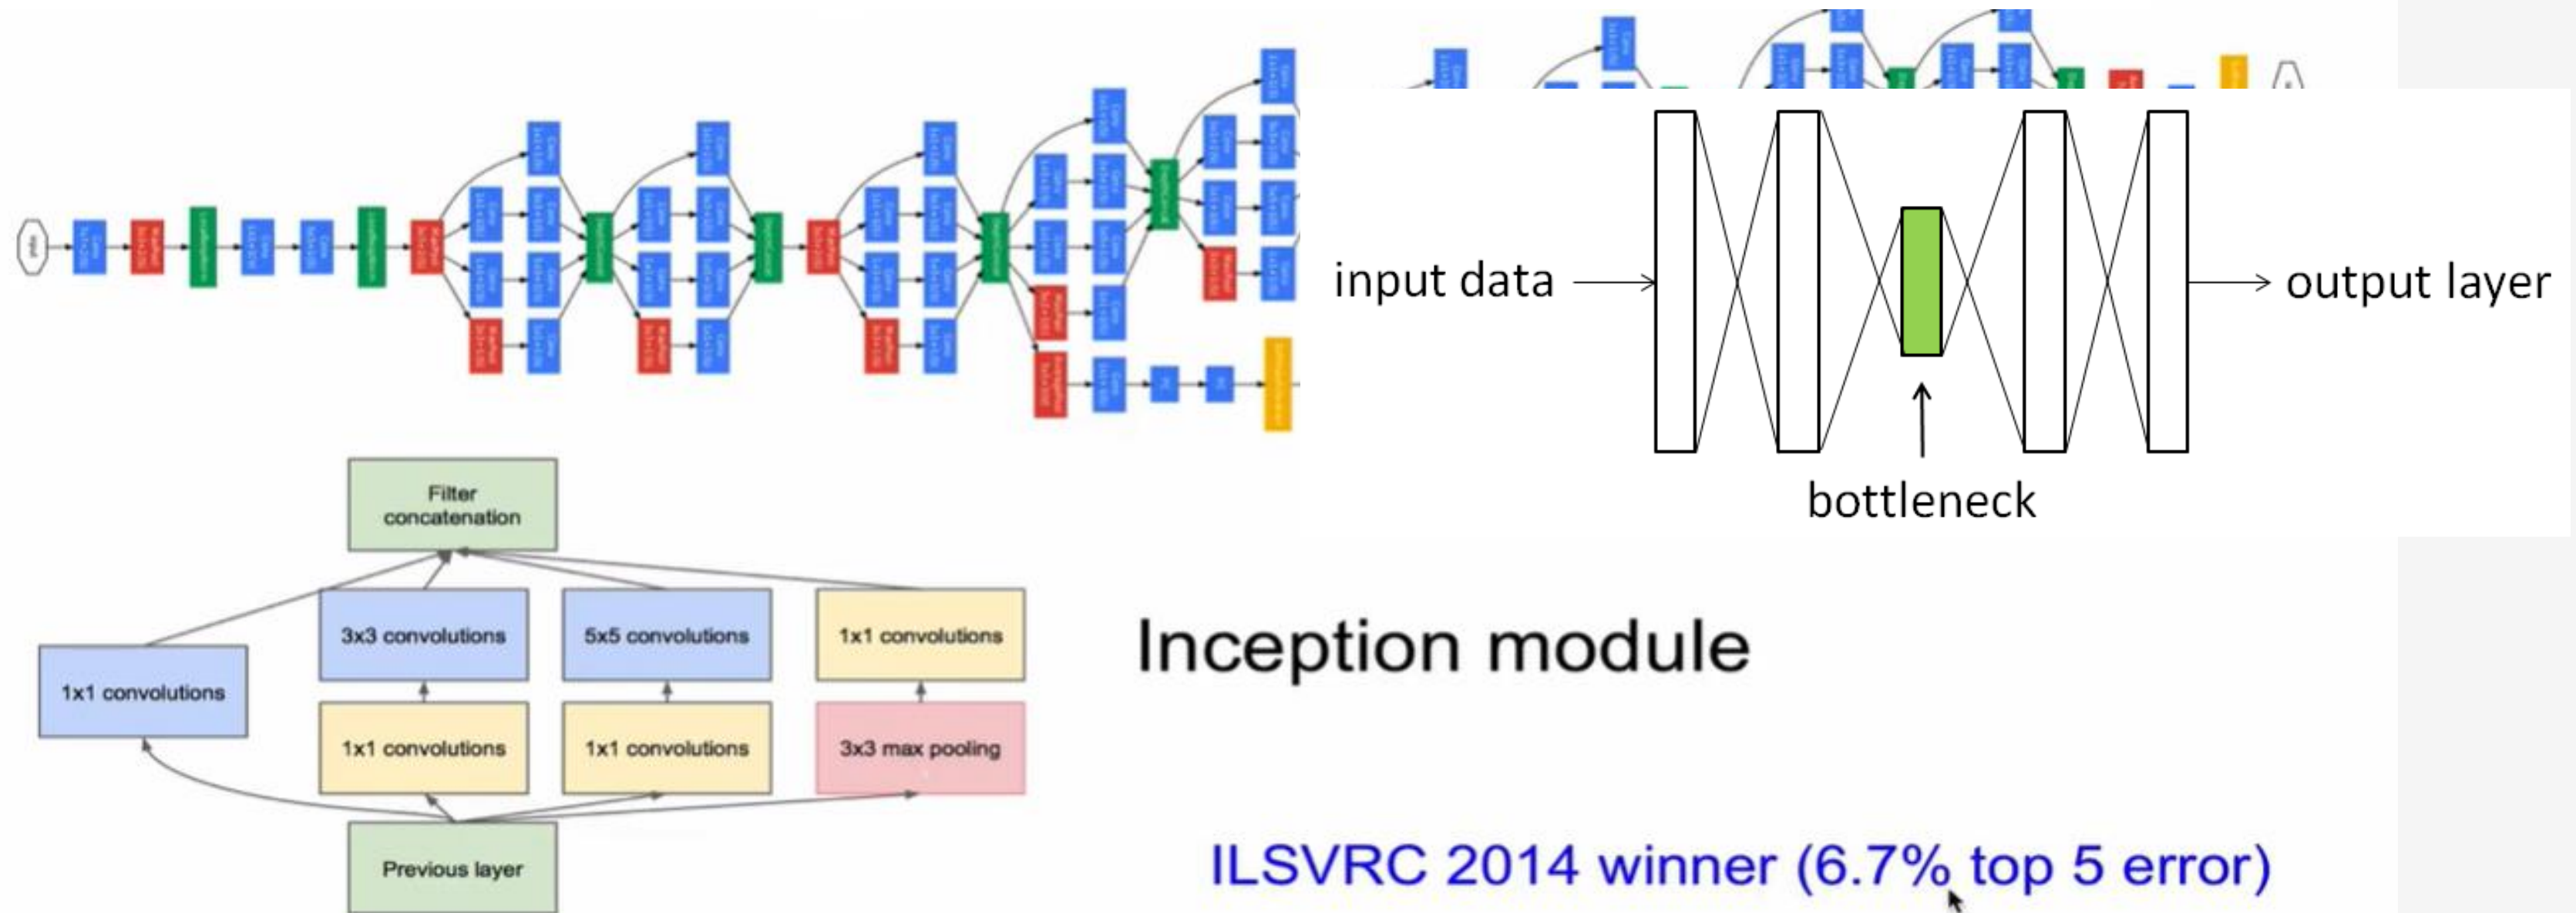
일정한 사이즈의 Conv filter (3*3) 중첩시켜 사용

-> 계산 비용 줄어듦, 성능 향상 O

-> 그러나, Gradient Vanishing 문제 발생

02. CNN의 흐름 및 모델

2014: GooLeNet – Going Deeper with Convolutions



02. CNN의 흐름 및 모델

Resea

MSRA @ ILSVRC & COCO 2015 Competitions

- **1st places** in all five main tracks

- ImageNet Classification: *"Ultra-deep"* (quote Yann) **152-layer** nets
- ImageNet Detection: **16%** better than 2nd
- ImageNet Localization: **27%** better than 2nd
- COCO Detection: **11%** better than 2nd
- COCO Segmentation: **12%** better than 2nd

Case Study: ResNet

[He et al., 2015]

- Batch Normalization after every CONV layer
- Xavier/2 initialization from He et al.
- SGD + Momentum (0.9)
- Learning rate: 0.1, divided by 10 when validation error plateaus
- Mini-batch size 256
- Weight decay of 1e-5
- No dropout used

02. CNN의 흐름 및 모델

2015: ResNet – Deep Residual Learning for Image Recognition


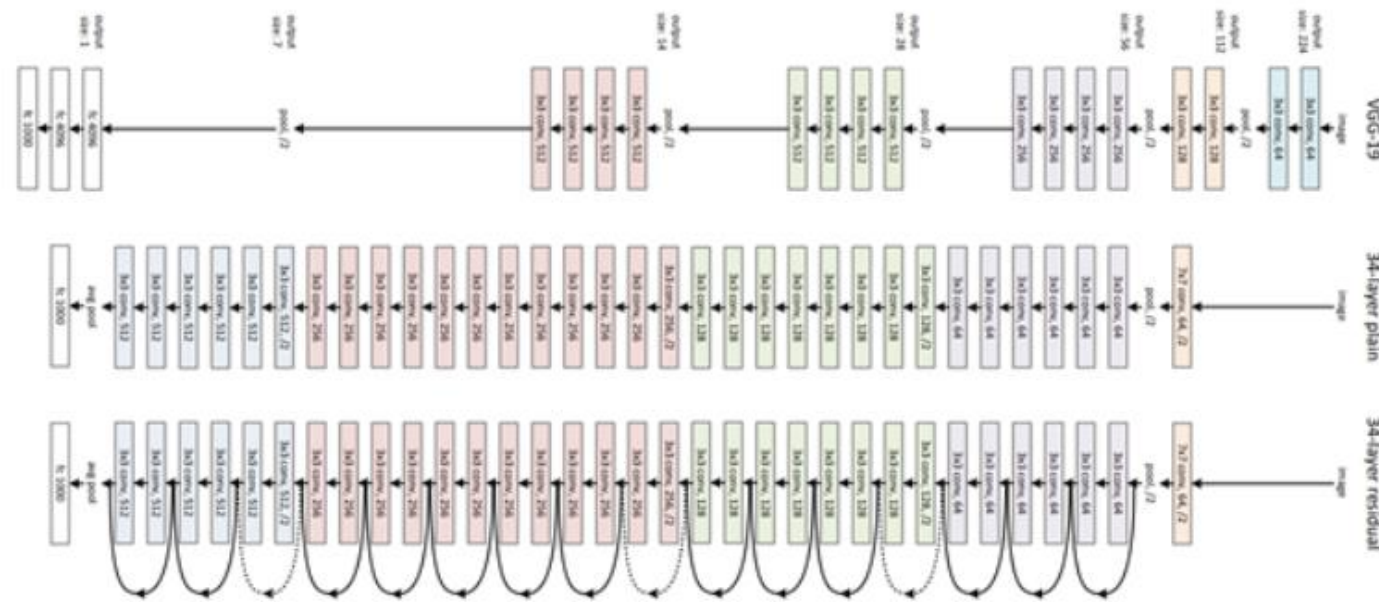
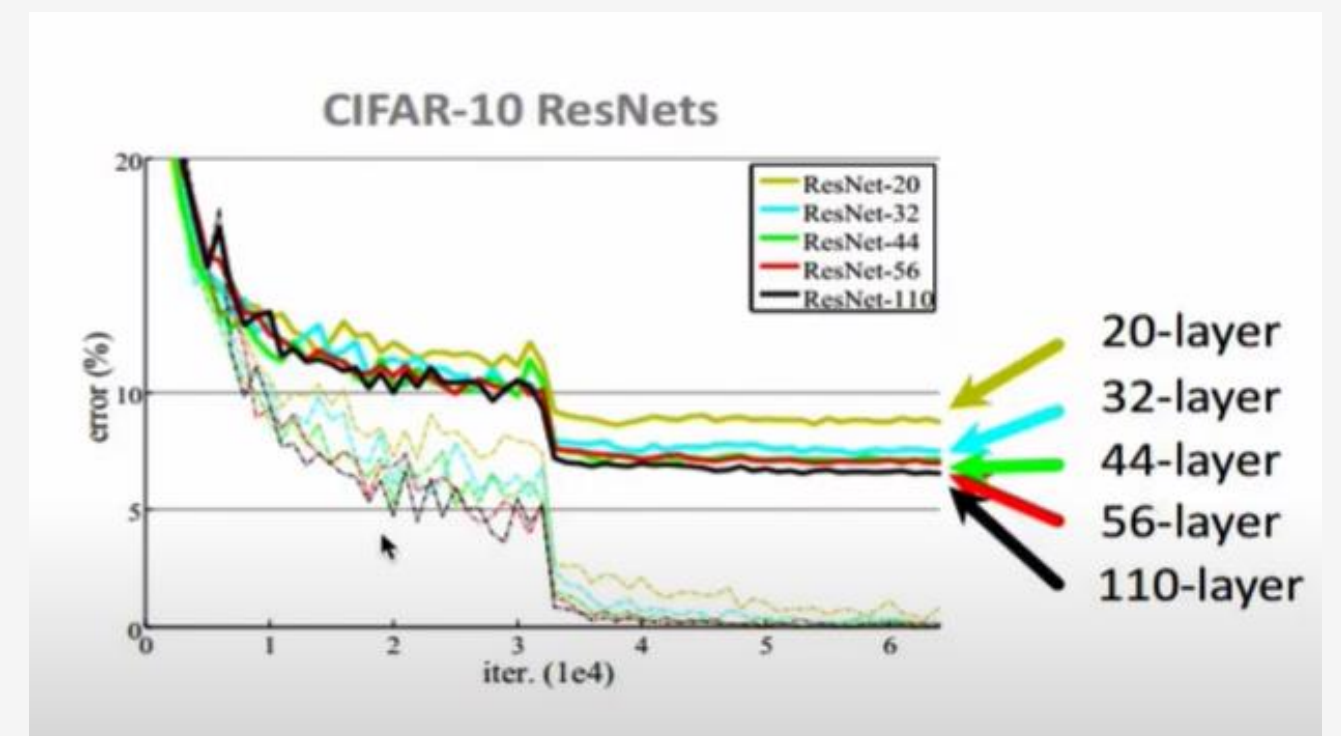




Figure 2. Residual learning: a building block.

Degradation problem

CIFAR-10 experiments




02. CNN의 흐름 및 모델



02. CNN의 흐름 및 모델

2017: DenseNet – Densely Connected Convolutional Networks



$$x_l = H_l([x_0, x_1, \dots, x_{l-1}]),$$

Figure 1: A 5-layer dense block w
Each layer takes all preceding feature

- 모든 Layer의 feature map 연결
-> Gradient vanishing 문제 완화
- 파라미터 수, 연산량이 적음.


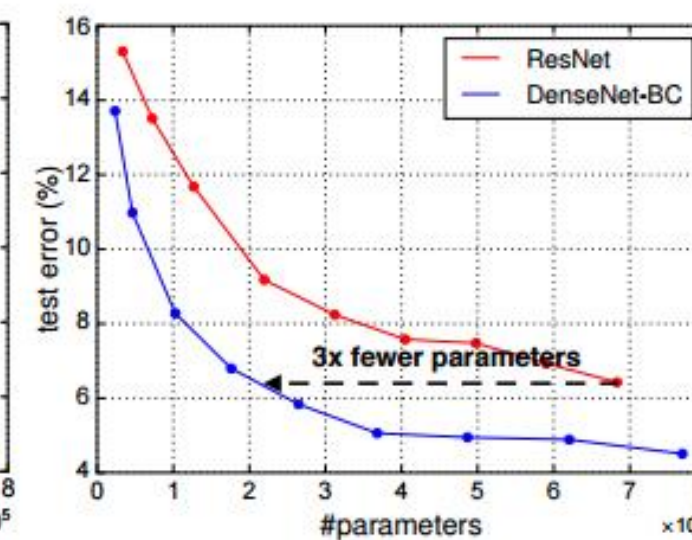



Figure 4: Left: Comparison of the parameter efficiency on C10+ between DenseNet variations. Middle: Comparison of the parameter efficiency between DenseNet-BC and (pre-activation) ResNets. DenseNet-BC requires about 1/3 of the parameters as ResNet to achieve comparable accuracy. Right: Training and testing curves of the 1001-layer pre-activation ResNet [12] with more than 10M parameters and a 100-layer DenseNet with only 0.8M parameters.

02. CNN의 흐름 및 모델

2017: MobileNet – MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Application

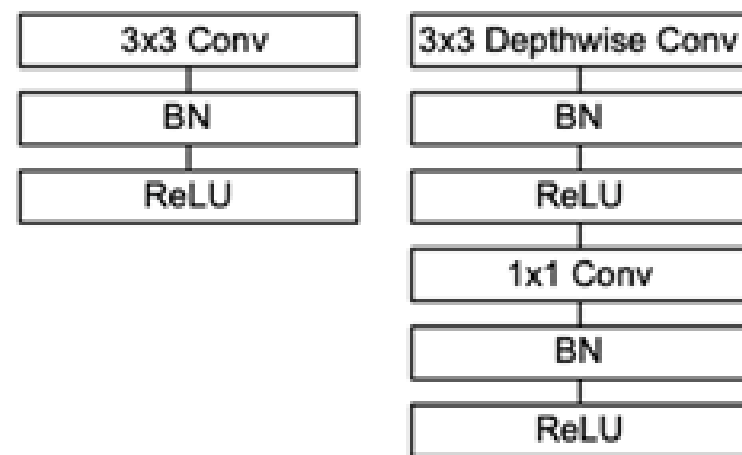


Figure 3. Left: Standard convolutional layer with batchnorm and ReLU. Right: Depthwise Separable convolutions with Depthwise and Pointwise layers followed by batchnorm and ReLU.

Table 1. MobileNet Body Architecture

Type / Stride	Filter Shape	Input Size
Conv / s2	$3 \times 3 \times 3 \times 32$	$224 \times 224 \times 3$
Conv dw / s1	$3 \times 3 \times 32 \text{ dw}$	$112 \times 112 \times 32$
Conv / s1	$1 \times 1 \times 32 \times 64$	$112 \times 112 \times 32$
Conv dw / s2	$3 \times 3 \times 64 \text{ dw}$	$112 \times 112 \times 64$
Conv / s1	$1 \times 1 \times 64 \times 128$	$56 \times 56 \times 64$
Conv dw / s1	$3 \times 3 \times 128 \text{ dw}$	$56 \times 56 \times 128$
Conv / s1	$1 \times 1 \times 128 \times 128$	$56 \times 56 \times 128$
Conv dw / s2	$3 \times 3 \times 128 \text{ dw}$	$56 \times 56 \times 128$
Conv / s1	$1 \times 1 \times 128 \times 256$	$28 \times 28 \times 128$
Conv dw / s1	$3 \times 3 \times 256 \text{ dw}$	$28 \times 28 \times 256$
Conv / s1	$1 \times 1 \times 256 \times 256$	$28 \times 28 \times 256$
Conv dw / s2	$3 \times 3 \times 256 \text{ dw}$	$28 \times 28 \times 256$
Conv / s1	$1 \times 1 \times 256 \times 512$	$14 \times 14 \times 256$
5×	Conv dw / s1	$3 \times 3 \times 512 \text{ dw}$
	Conv / s1	$1 \times 1 \times 512 \times 512$
	Conv dw / s2	$3 \times 3 \times 512 \text{ dw}$
	Conv / s1	$1 \times 1 \times 512 \times 1024$
	Conv dw / s2	$3 \times 3 \times 1024 \text{ dw}$
	Conv / s1	$1 \times 1 \times 1024 \times 1024$
	Avg Pool / s1	Pool 7×7
	FC / s1	1024×1000
	Softmax / s1	Classifier

2017: ShuffleNet – ShuffleNet: An Extremely Efficient Convolutional Neural Network for Mobile Devices

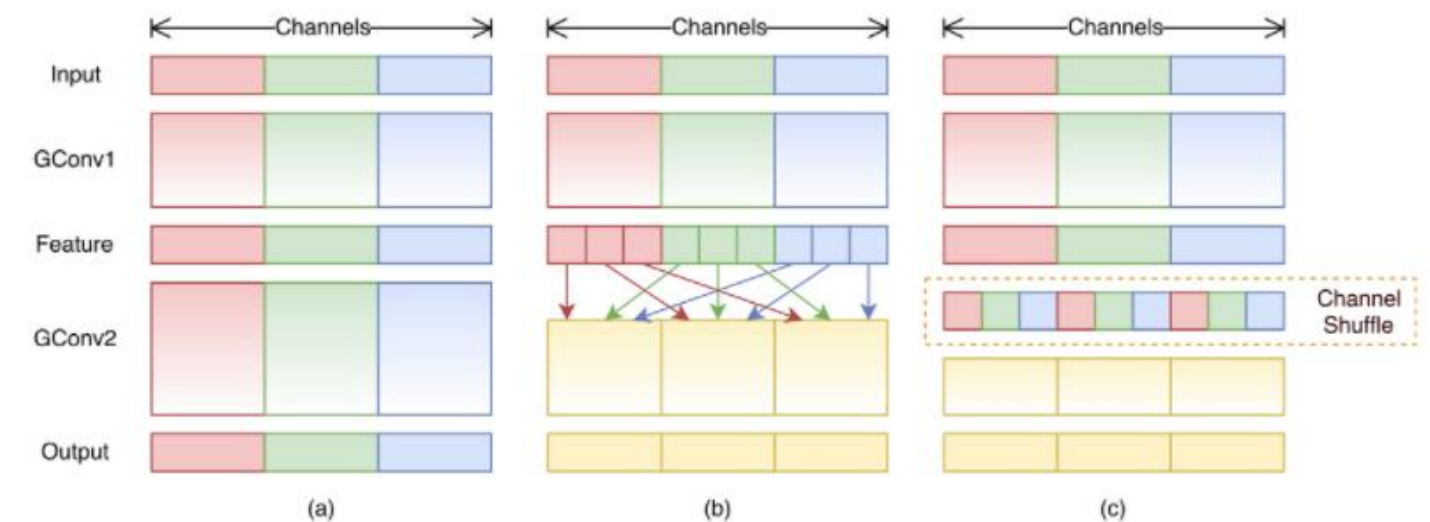


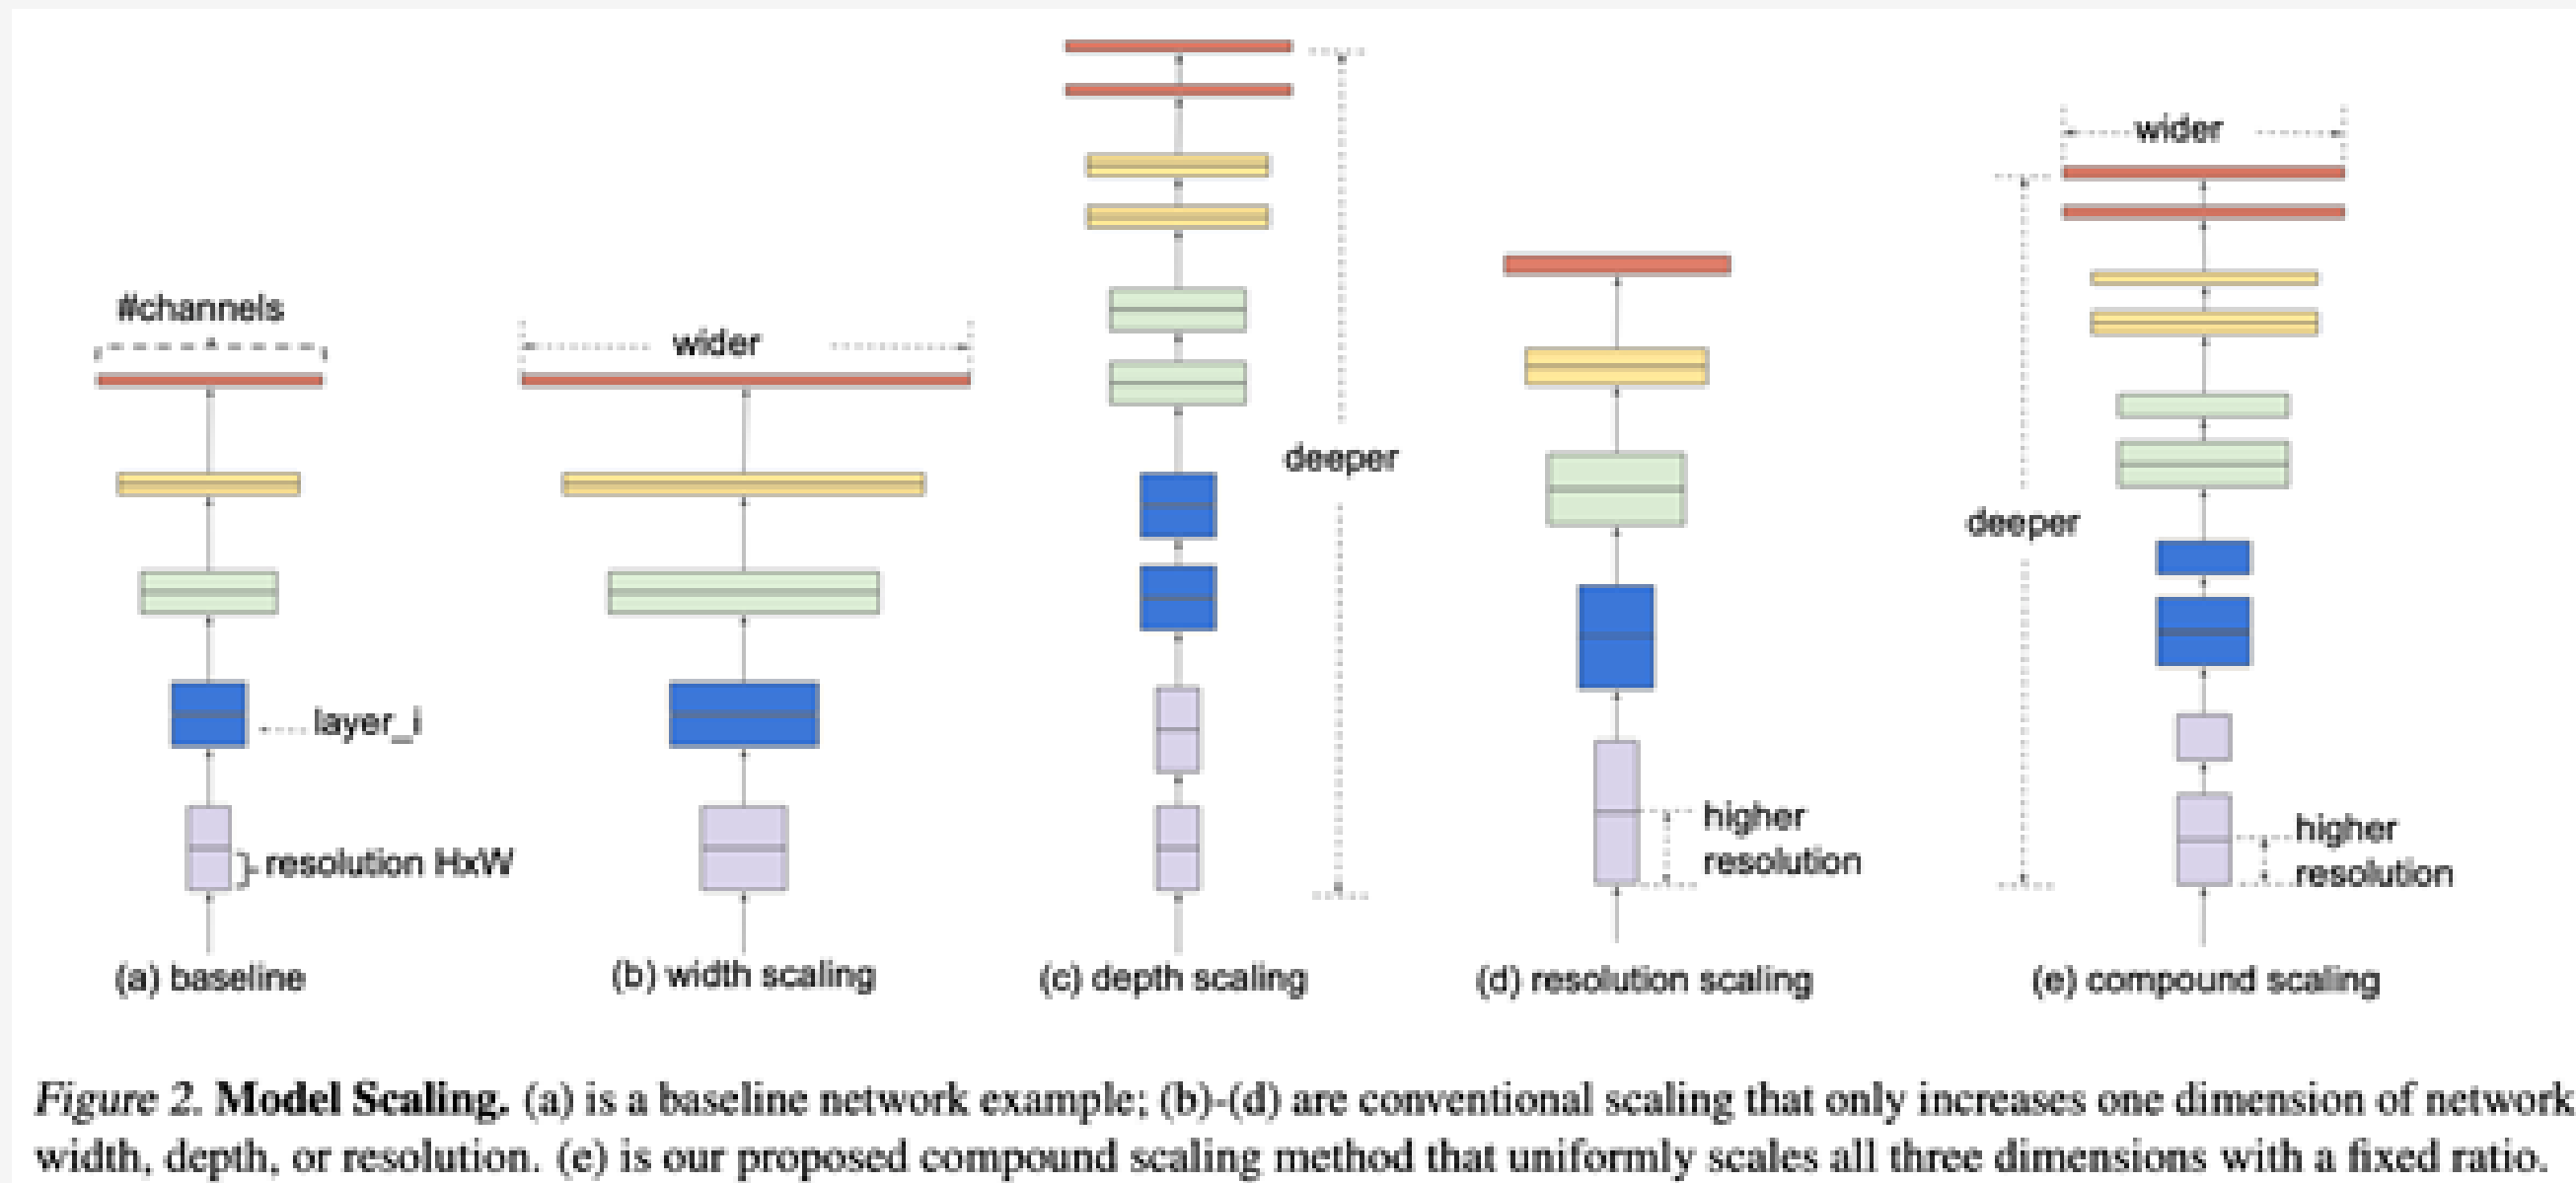
Figure 1. Channel shuffle with two stacked group convolutions. GConv stands for group convolution. a) two stacked convolution layers with the same number of groups. Each output channel only relates to the input channels within the group. No cross talk; b) input and output channels are fully related when GConv2 takes data from different groups after GConv1; c) an equivalent implementation to b) using channel shuffle.

02. CNN의 흐름 및 모델

2019: EfficientNet – EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks

- 모델의 크기를 어떻게 키울 것인가?

-> 모델 Scale up의 원리는 무엇일까?



02. CNN의 흐름 및 모델

○ Problem Formulation:

$$\mathcal{N} = \bigodot_{i=1 \dots s} \mathcal{F}_i^{L_i} (X_{\langle H_i, W_i, C_i \rangle})$$

- \mathcal{N} : CNN 모델
- \mathcal{F}_i : Height, Width, Channel로 구성된 **Layer** (Operator)
- L_i : Stage i 에서의 반복 수(Repeated Times in stage i)

$$\begin{aligned} \text{depth: } d &= \alpha^\phi \\ \text{width: } w &= \beta^\phi \\ \text{resolution: } r &= \gamma^\phi \\ \text{s.t. } \alpha \cdot \beta^2 \cdot \gamma^2 &\approx 2 \\ \alpha \geq 1, \beta \geq 1, \gamma &\geq 1 \end{aligned}$$




Figure 3. Scaling Up a Baseline Model with Different Network Width (w), Depth (d), and Resolution (r) Coefficients. Bigger networks with larger width, depth, or resolution tend to achieve higher accuracy, but the accuracy gain quickly saturate after reaching 80%, demonstrating the limitation of single dimension scaling. Baseline network is described in Table 1.

- Width :

- * 특징 : 1. 작은 모델의 사이즈를 키울 때 많이 사용됨
2. 세밀(Fine-Grained)한 Feature를 Capture 하기 위해 많이 사용됨
- * 단점 : 1. Shallow 모델 (= Layer가 낮은)에서, 상위 레벨의 복잡한 Feature를 파악하기 어려움
2. Width가 넓어짐에 따라 빠르게 포화(Saturate)됨

- Depth :

- * 장점 : 1. Capture Richer and more Complex Feature
- * 단점 : 1. Skip-Connection이나 Batch Norm으로 완화하고 있지만, Vanishing Gradient로 학습하기 어려움
2. 특정 층 이상으로 쌓이면, 성능이 저하됨

- Resolution :

- * 장점 : 1. 세밀(Fine-Grained)한 Feature를 Capture 하기 위해 많이 사용됨
2. Detection처럼 복잡한 Task에서 더 높은 Resolution을 사용
3. Resolution의 증가는 정확도에 영향을 줌
- * 단점 : 1. Resolution의 증가할수록, 성능 증가폭은 감소됨
($r = 1.0 - 224 \times 224$ / $r = 2.5 - 560 \times 560$)

02. CNN의 흐름 및 모델




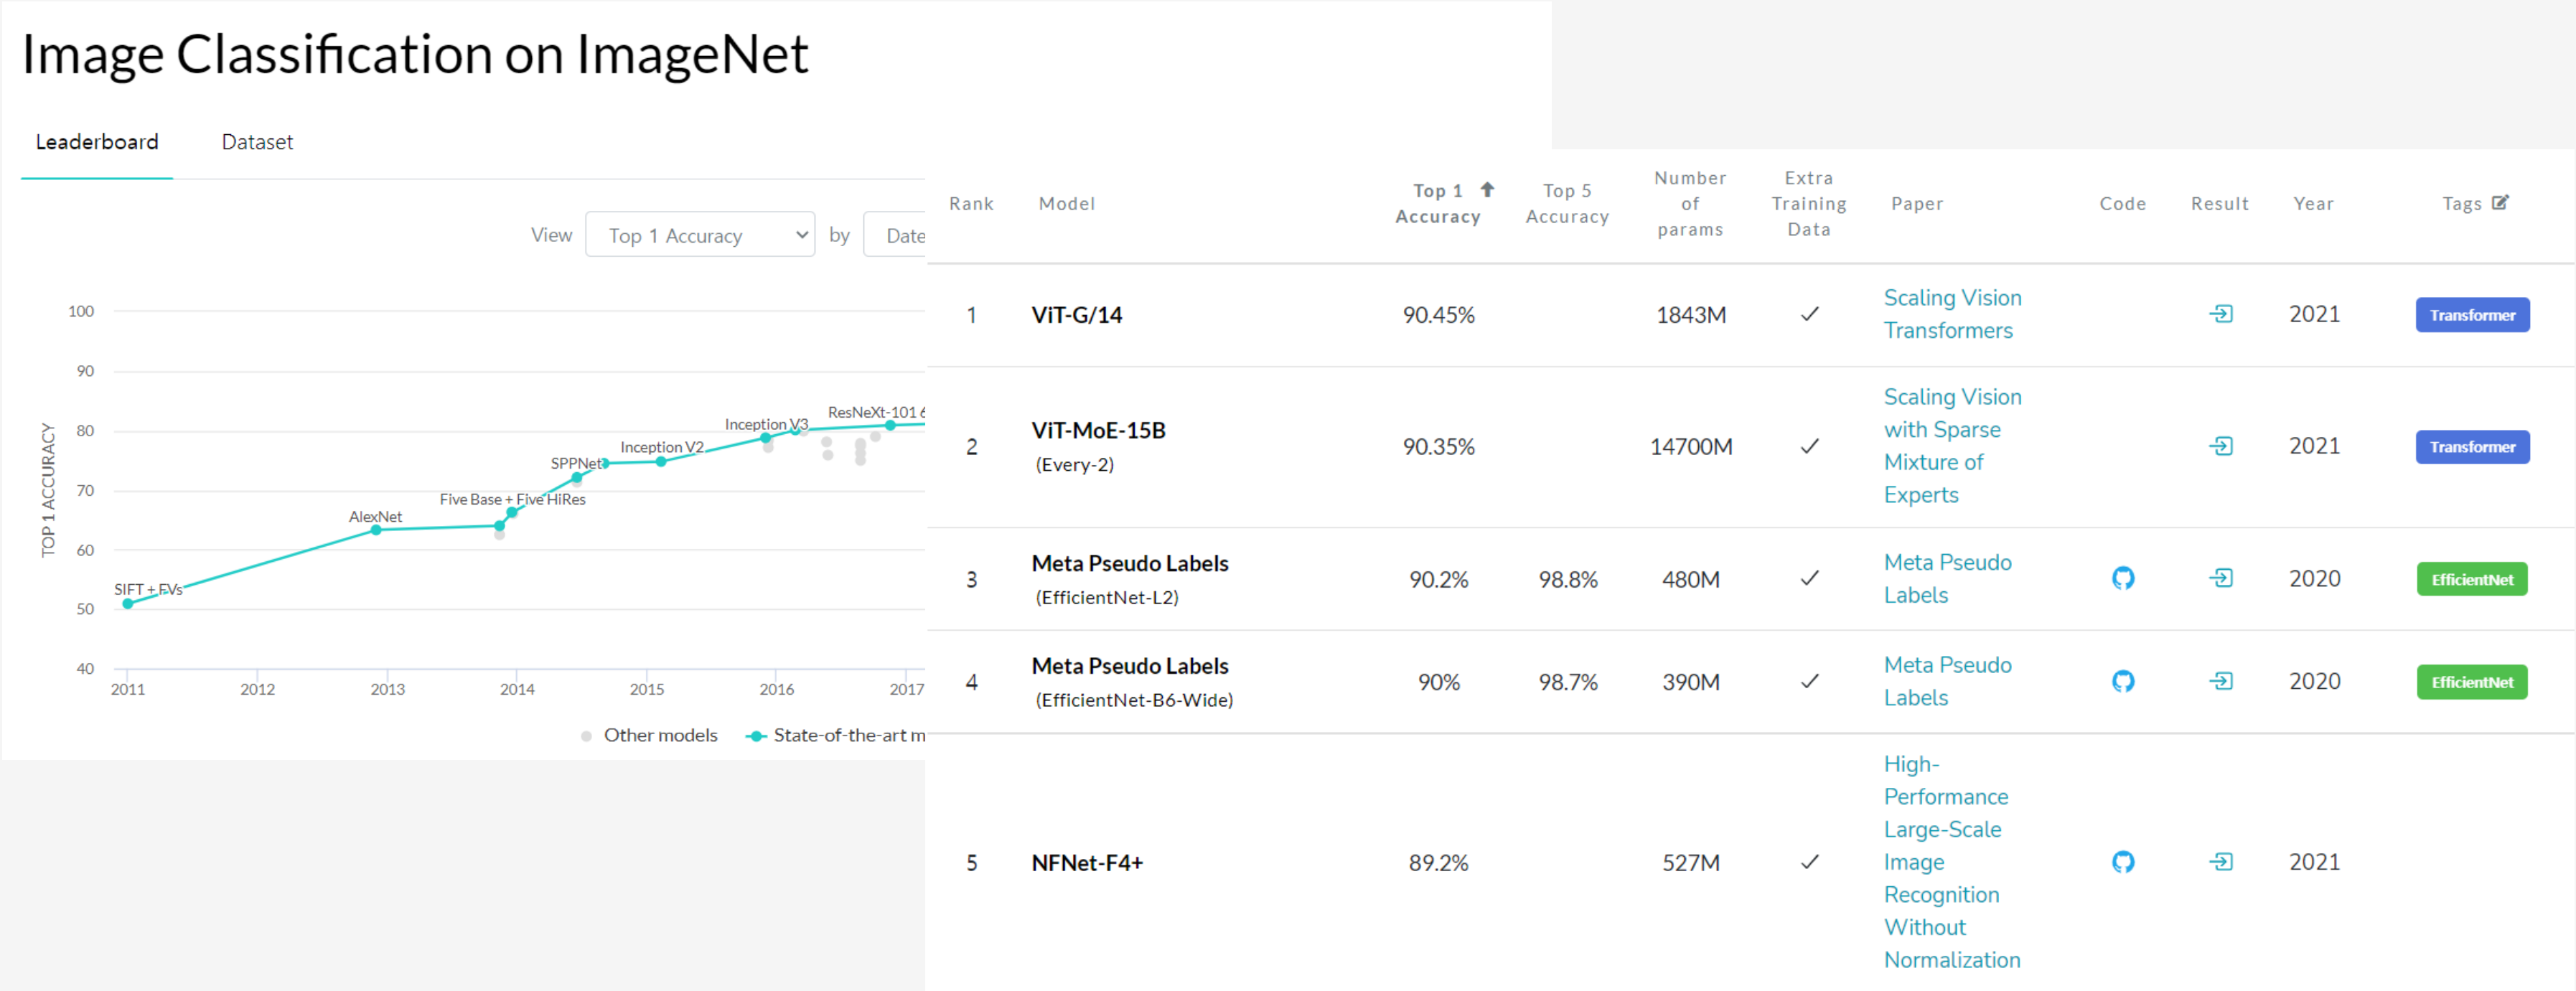
Figure 5. **FLOPS vs. ImageNet Accuracy** – Similar to Figure 1 except it compares FLOPS rather than model size.

Table 4. **Inference Latency Comparison** – Latency is measured with batch size 1 on a single core of Intel Xeon CPU E5-2690.

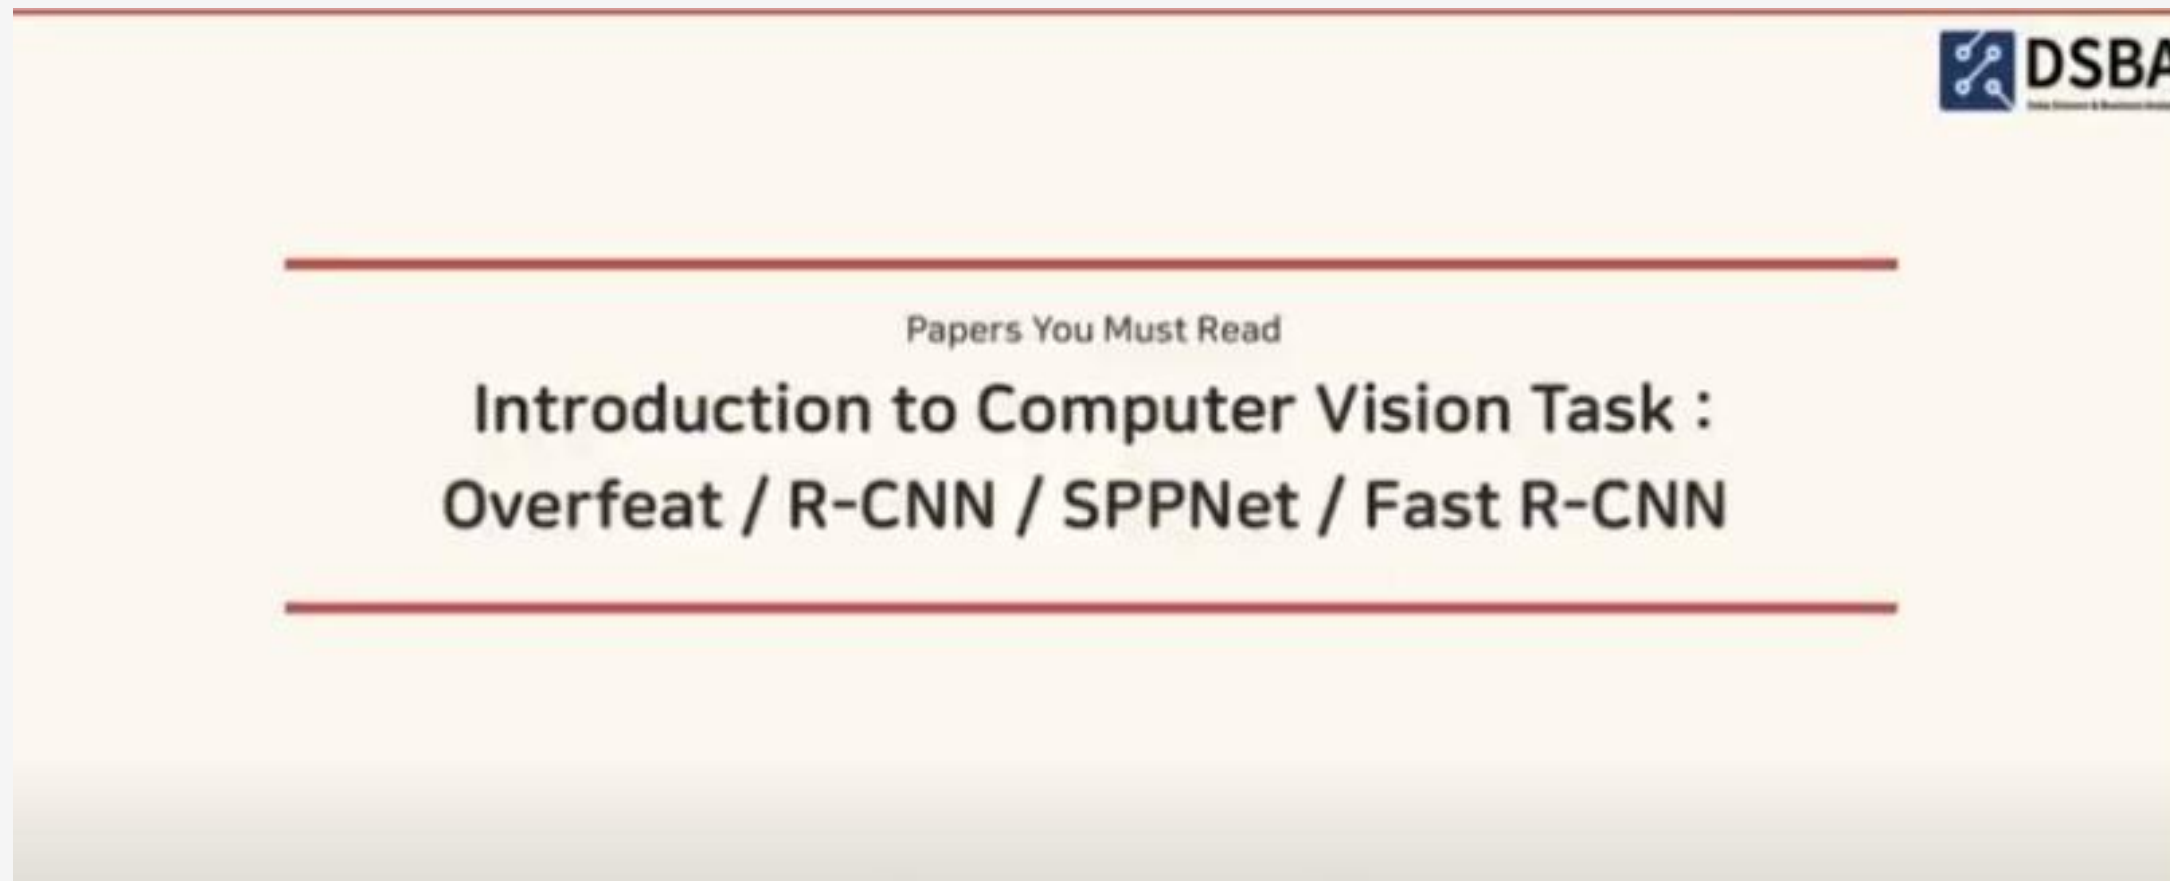
	Acc. @ Latency		Acc. @ Latency
ResNet-152	77.8% @ 0.554s	GPipe	84.3% @ 19.0s
EfficientNet-B1	78.8% @ 0.098s	EfficientNet-B7	84.4% @ 3.1s
Speedup	5.7x	Speedup	6.1x

02. CNN의 흐름 및 모델

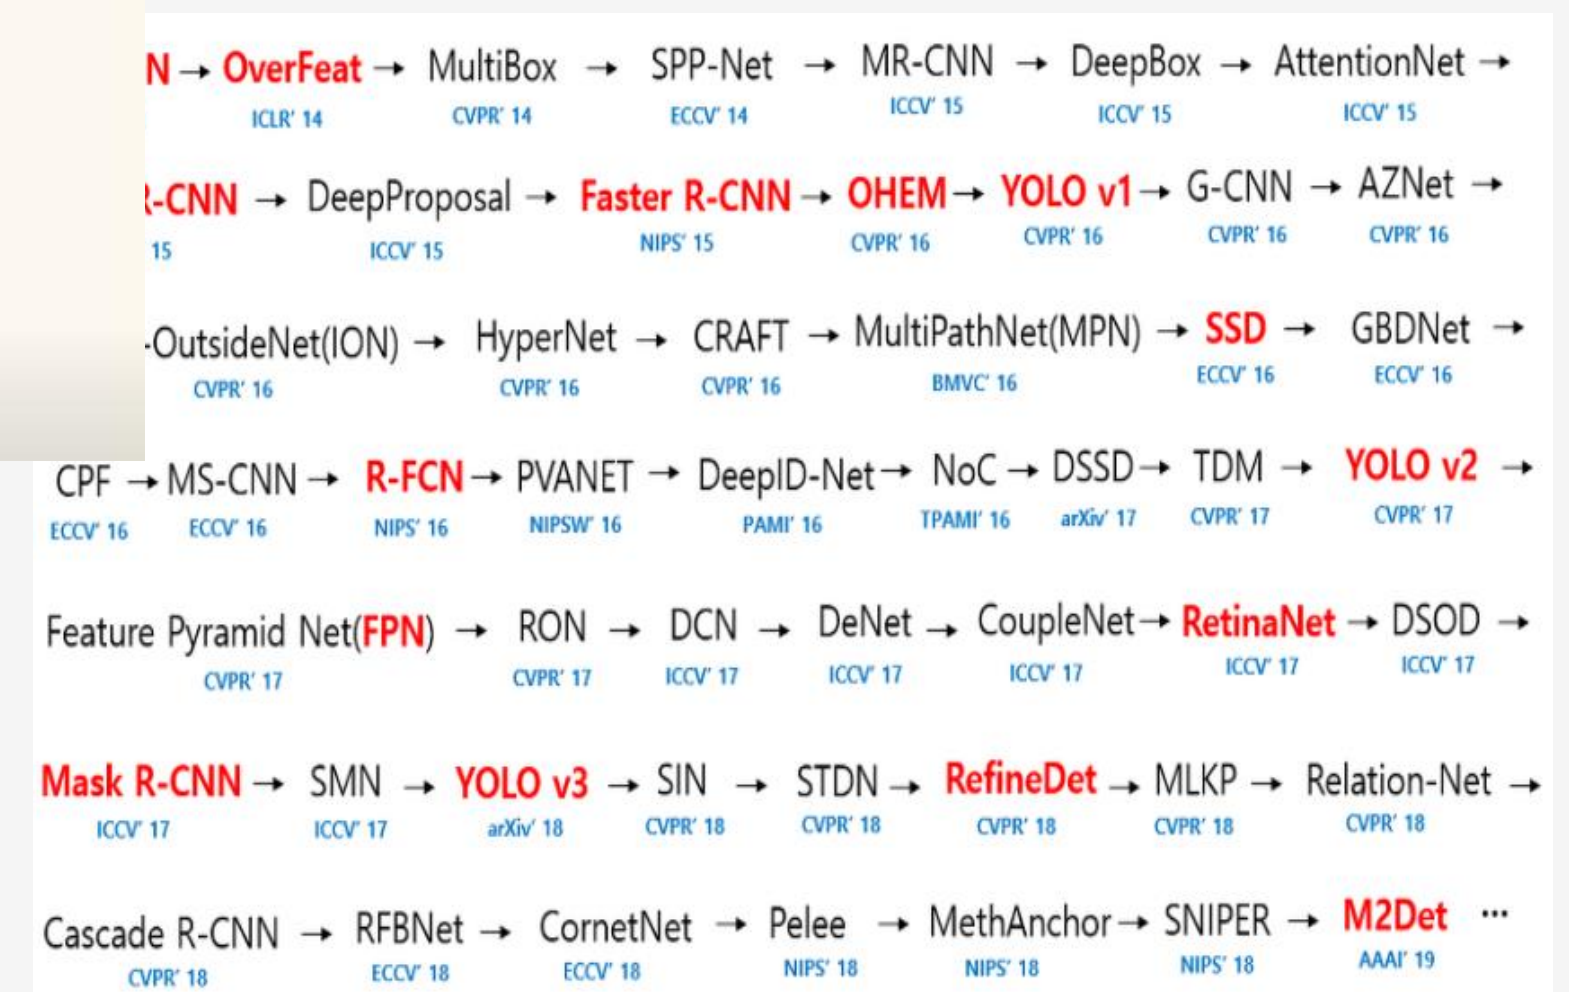
https://paperswithcode.com/sota/image-classification-on-imagenet



02. CNN의 흐름 및 모델



[\[Paper Review\] Introduction to Object Detection Task : Overfeat, RCNN, SPPNet, FastRCNN - YouTube](#)



THANK YOU

감사합니다.